

Elastic Search: Similarity & Scoring[1-4]

How Retrieved Documents are Ranked

Weimao Ke

Drexel University

Table of contents

1. Mapping and Built-in Similarity
2. Custom Similarities
3. Script Score Query
4. Practical Scoring with Relevance Signals

Mapping and Built-in Similarity

Mapping Similarity

With Mapping:

- You can configure a scoring or similarity function for each field.
- Basic similarity functions include: *BM25* (default), *classic* (TF*IDF), and *boolean*
- Custom similarities can be configured with related parameters.

This changes the way in which retrieved documents are ranked (sorted).

It is the most important function of a search engine.

Mapping Similarity

Example:

```
1  PUT my_index
2  {
3    "mappings": {
4      "properties": {
5        "field1": {
6          "type": "text"
7        },
8        "field2": {
9          "type": "text",
10         "similarity": "boolean"
11       }
12     }
13   }
14 }
```

- By default, *field1* uses *BM25* similarity.
- *Field2* is configured to use the *boolean* similarity.

Mapping Similarity

You can change the default similarity for a new index:

```
1  PUT /index
2  {
3    "settings": {
4      "index": {
5        "similarity": {
6          "default": {
7            "type": "boolean"
8          }
9        }
10     }
11  }
12 }
```

This is done when the index is created.

Mapping Similarity

You can change the default similarity of an existing index.

You will have to close the index to change it:

```
1  POST /index/_close
2
3  PUT /index/_settings
4  {
5    "index": {
6      "similarity": {
7        "default": {
8          "type": "boolean"
9        }
10     }
11  }
12 }
```

And reopen it after the change:

```
1  POST /index/_open
```


Custom Similarities

Similarity Modules

1

BM25 similarity (default):

- Regarded as a version of $TF \cdot IDF$, with TF normalization.
- k_1 : controls TF normalization (saturation), default 1.2;
- b : controls document length normalization, default 0.75.



DFR (Divergence from Randomness) similarity:

- *basic_model*: possible values *g*, *if*, *in*, and *ine*;
- *after_effect*: possible values *b* and *l*;
- *normalization*: possible values *no*, *h1*, *h2*, *h3*, and *z*.

All options but the first option need a normalization value.

Other similarity functions include:

- DFI (Divergence from Independence) similarity
- IB (information based model) similarity
- LM Dirichlet similarity
- LM Jelinek Mercer similarity

Mapping Similarity

One can create a customized version of a built-in similarity function:

```
1  PUT /index
2  {
3    "settings": {
4      "index": {
5        "similarity": {
6          "my_similarity": {
7            "type": "DFR",
8            "basic_model": "g",
9            "after_effect": "l",
10           "normalization": "h2",
11           "normalization.h2.c": "3.0"
12         }
13       }
14     }
15   }
16 }
```

Mapping Similarity

And then use the custom version with mapping:

```
1  PUT /index/_mapping
2  {
3    "properties" : {
4      "title" : { "type" : "text", "similarity" : "
                    my_similarity" }
5    }
6  }
```

Mapping Similarity

ElasticSearch also allows scripted similarity:

```
1  PUT /index
2  {
3    "settings": {
4      "number_of_shards": 1,
5      "similarity": {
6        "custom_tfidf": {
7          "type": "scripted",
8          "script": {
9            "source": "
10             double tf = Math.sqrt(doc.freq);
11             double idf = Math.log((field.docCount+1.0)/(term.
12               docFreq+1.0)) + 1.0;
13             double norm = 1/Math.sqrt(doc.length);
14             return query.boost * tf * idf * norm;"
15           }
16         }
17       },
```

Mapping Similarity

Scripted similarity (continued):

```
1  "mappings": {  
2    "properties": {  
3      "field": {  
4        "type": "text",  
5        "similarity": "custom_tfidf"  
6      }  
7    }  
8  }  
9  }
```

Mapping Similarity

Note:

- *doc.freq* is NOT document frequency (DF), but a term's frequency in the document (TF).
- *term.docFreq* is document frequency (DF).

Rules on scripted similarity:

- Returned scores must be positive;
- All other variables remaining equal, scores must not decrease when term frequencies (TF) increases;
- All other variables remaining equal, scores must not increase when document length increases.

Mapping Similarity

You may also use *weight_script* for more efficient scoring:

```
1 PUT /index
2 {
3   "settings": {
4     "number_of_shards": 1,
5     "similarity": {
6       "scripted_tfidf": {
7         "type": "scripted",
8         "weight_script": {
9           "source": "double idf =
10             Math.log((field.docCount+1.0)/(term.docFreq+1.0)) + 1.0;
11             return query.boost * idf;"
12         },
13         "script": {
14           "source": "
15             double tf = Math.sqrt(doc.freq);
16             double norm = 1/Math.sqrt(doc.length);
17             return weight * tf * norm;"
18         }
19       }
20     }
21   },
```

Mapping Similarity

weight_script (continued):

```
1  "mappings": {  
2    "properties": {  
3      "field": {  
4        "type": "text",  
5        "similarity": "scripted_tfidf"  
6      }  
7    }  
8  }  
9  }
```

Script Score Query

Query with a score script:

- *script_score* provides a custom score as part of a query request.
- Changes the way in which returned documents are scored and ranked for the query.

Script Score Query

Example script score query:

```
1  GET /_search
2  {
3      "query" : {
4          "script_score" : {
5              "query" : {
6                  "match" : { "message": "Best Video Tutorials"
7                      }
8              },
9              "script" : {
10                 "source" : "doc['likes'].value / 10 "
11             }
12         }
13     }
```

Suppose `doc['likes']` is a field with the count of likes.

Here it ranks documents according to the number of likes.

Parameters for *script_score*:

- *query*: the query to search and return documents;
- *script*: the script used to compute the score of each returned document;
- *min_score* (optional): the threshold score for a document to be included in the search results (hits).

Note that *_score* variable has the current relevance score for each document and can be used as part of the new scoring function.

Predefined functions for *script_score*

Saturation:

$$\text{saturation}(\text{value}, k) = \text{value} / (k + \text{value})$$

```
1  "script" : {  
2      "source" : "saturation(doc[ 'likes ' ].value , 1)"  
3  }
```

that is $\frac{\text{likes}}{1+\text{likes}}$:

- 0 when *likes* = 0;
- 0.5 when *likes* = 1;
- 1 when *likes* $\rightarrow \infty$

Predefined functions for *script_score*

Sigmoid:

$$\text{sigmoid}(\text{value}, k, a) = \text{value}^a / (k^a + \text{value}^a)$$

```
1  "script" : {  
2      "source" : "sigmoid(doc['likes'].value, 2, 1)"  
3  }
```

that is $\frac{\text{likes}^1}{2^1 + \text{likes}^1} = \frac{\text{likes}}{2 + \text{likes}}$

Predefined functions for *script_score*

Random score:

randomScore(< seed >, < fieldName >)

```
1  "script" : {  
2    "source" : "randomScore(100)"  
3  }
```

- This generates some randomness (noise) in the ranking, e.g. for random sampling in the search results.
- < *fieldName* > parameter, as the source of randomness, is optional. Document ID is used by default to generate the random score.

Predefined functions for *script_score*

Decay functions:

- Provides scoring based on some kind of distance for numeric values, dates, and geo locations.

An example for geo fields:

```
1  "script" : {
2      "source" : "decayGeoExp(params.origin, params.scale,
3          params.offset, params.decay, doc['location'].value)",
4      "params": {
5          "origin": "40, -70.12",
6          "scale": "200km",
7          "offset": "0km",
8          "decay" : 0.2
9      }
10 }
```

Predefined functions for *script_score*

Functions for vector fields (experimental):

- Functions to be used on *dense_vector* and *sparse_vector* fields.
- A score is computed for every matched document so try to limit the number of docs with a query parameter.

An index with vector fields:

```
1  PUT my_index
2  {
3    "mappings": {
4      "properties": {
5        "my_dense_vector": {
6          "type": "dense_vector",
7          "dims": 3
8        },
9        "my_sparse_vector" : {
10         "type" : "sparse_vector"
11       },
12       "status" : {
13         "type" : "keyword"
14       }
15     }
16  }
```

Predefined functions for *script_score*

Example vector data:

```
1  PUT my_index/_doc/1
2  {
3    "my_dense_vector": [0.5, 10, 6],
4    "my_sparse_vector": {"2": 1.5, "15" : 2, "50": -1.1, "4545": 1.1},
5    "status" : "published"
6  }
7
8  PUT my_index/_doc/2
9  {
10   "my_dense_vector": [-0.5, 10, 10],
11   "my_sparse_vector": {"2": 2.5, "10" : 1.3, "55": -2.3, "113": 1.6},
12   "status" : "published"
13 }
```

Predefined functions for *script_score*

A script score query using cosine similarity:

```
1  GET my_index/_search
2  {
3    "query": {
4      "script_score": {
5        "query" : {
6          "bool" : {
7            "filter" : {
8              "term" : {
9                "status" : "published"
10             }
11          }
12        }
13      },
14      "script": {
15        "source": "cosineSimilarity(params.query_vector, doc['
16                  my_dense_vector']) + 1.0",
17        "params": {
18          "query_vector": [4, 3.4, -0.2]
19        }
20      }
21    }
22  }
```

Predefined functions for *script_score*

Cosine similarity of sparse vectors:

```
1  GET my_index/_search
2  {
3    "query": {
4      "script_score": {
5        "query" : {
6          "bool" : {
7            "filter" : {
8              "term" : {
9                "status" : "published"
10             }
11          }
12        }
13      },
14      "script": {
15        "source": "cosineSimilaritySparse(params.query_vector, doc['
16                  my_sparse_vector']) + 1.0",
17        "params": {
18          "query_vector": {"2": 0.5, "10" : 111.3, "50": -1.3, "113":
19                          14.8, "4545": 156.0}
20        }
21      }
22    }
```

Other functions for vector fields:

- *dotProduct* computes the dot product of two vectors;
- *dotProductSparse* is the sparse vector version for dot product calculation;
- *l1norm* or *l1normSparse* calculates the L^1 distance, i.e. Manhattan distance;
- *l2norm* or *l2normSparse* calculates the L^2 distance, i.e. Euclidean distance

Practical Scoring with Relevance Signals

Additional relevance signals:

- Relevance based on terms is not the only source of evidence.
- Other signals correlated with relevance: quality, reputation/popularity, importance, etc.
- Hyperlinks, PageRanks, # of likes, etc. are signals that can be used in relevance scoring (ranking).
- Can be implemented using *script_score* or *rank_feature* queries.

Relevance Signals

Scoring with PageRank (using *script_score*):

```
1  GET index/_search
2  {
3      "query" : {
4          "script_score" : {
5              "query" : {
6                  "match" : { "body": "University" }
7              },
8              "script" : {
9                  "source" :
10                     "_score * saturation(doc['pagerank'].value, 10)"
11              }
12          }
13      }
14  }
```

that is $bm25_score \times \frac{pagerank}{10 + pagerank}$

Here *pagerank* must be mapped as a *numeric* field.

Relevance Signals

Scoring with PageRank (using *rank_feature*):

```
1  GET _search
2  {
3      "query" : {
4          "bool" : {
5              "must": {
6                  "match": { "body": "University" }
7              },
8              "should": {
9                  "rank_feature": {
10                     "field": "pagerank",
11                     "saturation": {
12                         "pivot": 10
13                     }
14                 }
15             }
16         }
17     }
18 }
```

Here *pagerank* must be mapped as a *rank_feature* field.

Comparing *script_score* vs. *rank_feature*:

- *script_score* is more flexible and offers more ways to combine text relevance with other signals.
- *rank_feature* is limited but more efficient (faster);
- *rank_feature* fields are indexed signals, based on which the query can skip non-competitive documents and find the top matches much faster.

More on *rank_feature*:

- *rank_feature* is typically used in the *should* (OR) clause of a *bool* query;
- this is to *add* the signal to the other scores.
- supports a few functions: saturation, logarithm, and sigmoid.
- must setup *rank_feature* or *rank_features* fields with the index.

Relevance Signals

Setup an index for *rank_feature*:

```
1  PUT /test
2  {
3    "mappings": {
4      "properties": {
5        "pagerank": {
6          "type": "rank_feature"
7        },
8        "url_length": {
9          "type": "rank_feature",
10         "positive_score_impact": false
11       },
12       "topics": {
13         "type": "rank_features"
14       }
15     }
16   }
17 }
```

positive_score_impact : *false* for URL's negative impact on relevance; *rank_features* for the topics field that will contain a list of topics and their relevance.

Relevance Signals

Index documents with *rank_feature(s)*:

```
1  PUT /test/_doc/1?refresh
2  {
3    "url": "http://en.wikipedia.org/wiki/2016_Summer_Olympics",
4    "content": "Rio 2016",
5    "pagerank": 75.3,
6    "url_length": 42,
7    "topics": {
8      "sports": 50,
9      "brazil": 30
10   }
11 }
12
13 PUT /test/_doc/2?refresh
14 {
15   "url": "http://en.wikipedia.org/wiki/2016_Brazilian_Grand_Prix",
16   "content": "Formula One motor race held on 13 November 2016",
17   "pagerank": 38.9,
18   "url_length": 47,
19   "topics": {
20     "sports": 35,
21     "formula one": 65,
22     "brazil": 20
23   }
24 }
```

Relevance Signals

Query with *rank_feature(s)*:

```
1 GET /test/_search
2 {
3   "query": {
4     "bool": {
5       "must": [
6         {
7           "match": {
8             "content": "2016"
9           }
10        }
11      ],
12      "should": [
13        {
14          "rank_feature": {
15            "field": "pagerank"
16          }
17        },
```


Relevance Signals

Query with *rank_feature(s)* (continued):

```
1      {
2          "rank_feature": {
3              "field": "url_length",
4              "boost": 0.1
5          }
6      },
7      {
8          "rank_feature": {
9              "field": "topics.sports",
10             "boost": 0.4
11         }
12     }
13 ]
14 }
15 }
16 }
```

boost (relative to 1.0) is to boost – increase (> 1) or decrease (< 1) – a relevance score.

Relevance Signals

Saturation – example $\frac{\text{pagerank}}{\text{pagerank}+8}$:

```
1 GET /test/_search
2 {
3   "query": {
4     "rank_feature": {
5       "field": "pagerank",
6       "saturation": {
7         "pivot": 8
8       }
9     }
10  }
11 }
```

Relevance Signals

Logarithm – example $\log 4 + \text{pagerank}$:

```
1 GET /test/_search
2 {
3   "query": {
4     "rank_feature": {
5       "field": "pagerank",
6       "log": {
7         "scaling_factor": 4
8       }
9     }
10  }
11 }
```

Relevance Signals

Sigmoid – example $\frac{\text{pagerank}^{0.6}}{\text{pagerank}^{0.6} + 7^{0.6}}$:

```
1 GET /test/_search
2 {
3   "query": {
4     "rank_feature": {
5       "field": "pagerank",
6       "sigmoid": {
7         "pivot": 7,
8         "exponent": 0.6
9       }
10    }
11  }
12 }
```

References

- [1] elastic.co. Elasticsearch reference [7.5]: Index modules: Similarity.
<https://www.elastic.co/guide/en/elasticsearch/reference/current/index-modules-similarity.html>, .
Accessed: 2020-1-16.
- [2] elastic.co. Elasticsearch reference [7.5]: How to: Recipes: Incorporating static relevance signals into the score.
<https://www.elastic.co/guide/en/elasticsearch/reference/current/static-scoring-signals.html>, .
Accessed: 2020-1-16.

- [3] elastic.co. Elasticsearch reference [7.5]: Query dsl: Script score query. <https://www.elastic.co/guide/en/elasticsearch/reference/current/query-dsl-script-score-query.html>, . Accessed: 2020-1-16.
- [4] elastic.co. Elasticsearch reference [7.5]: Mapping: Similarity. <https://www.elastic.co/guide/en/elasticsearch/reference/current/similarity.html>, . Accessed: 2020-1-16.